

Computer vision and lane detection

Nacho Sañudo

University of Modena and Reggio Emilia

Ignacio.sanudoolmedo@unimore.it



UNIMORE
UNIVERSITÀ DEGLI STUDI DI
MODENA E REGGIO EMILIA

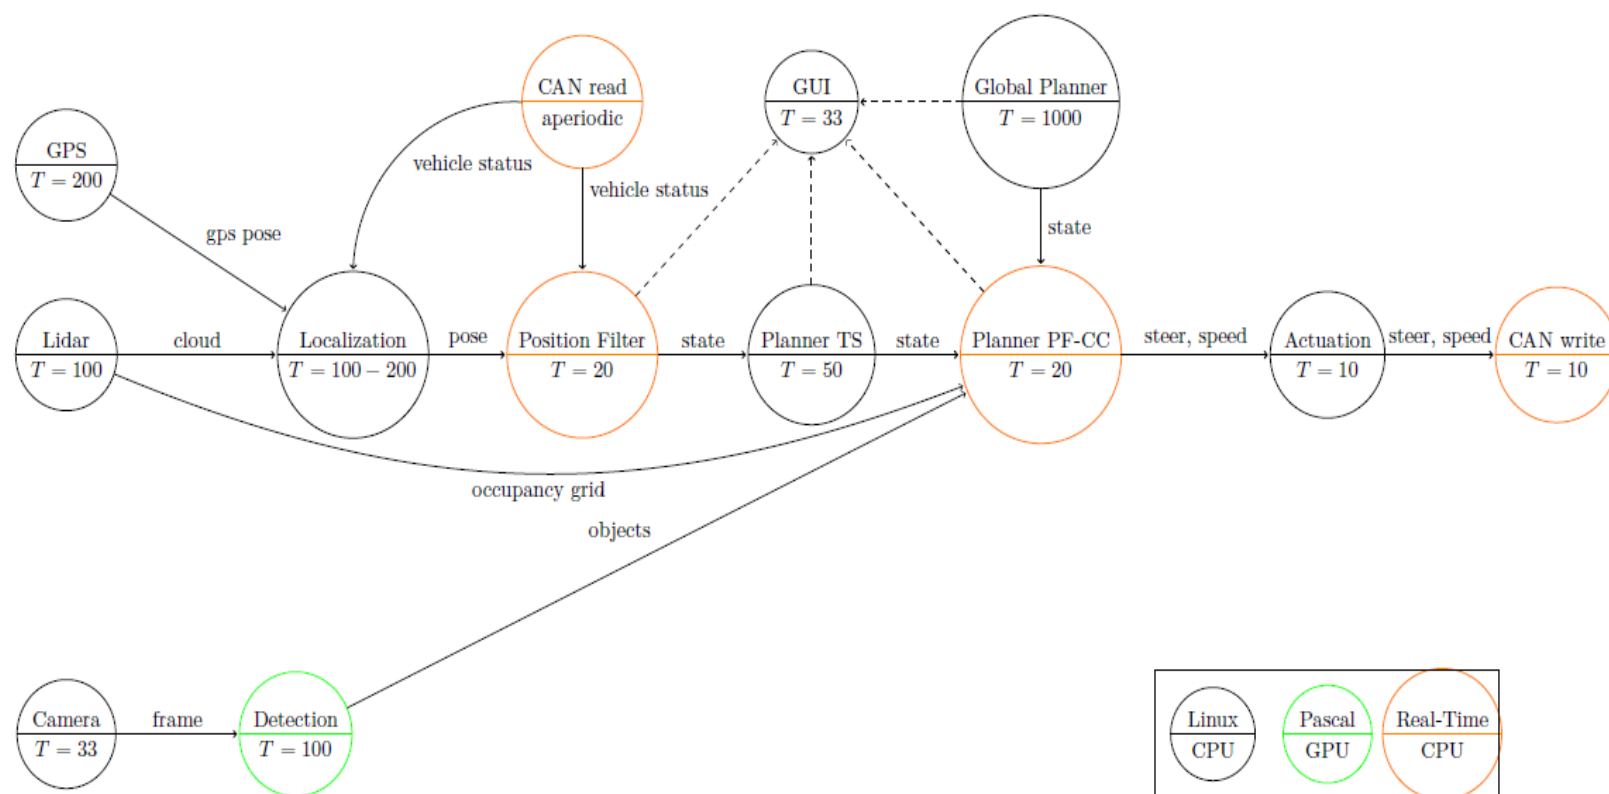
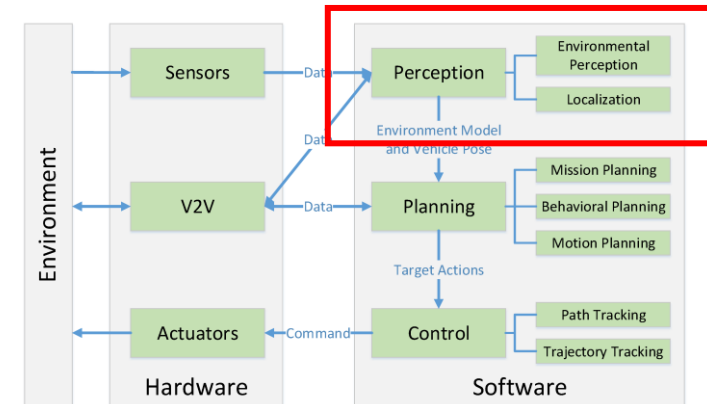


Contents

- › Perception
- › Cameras
 - Images, RGB model
- › Computer vision
 - Machine learning
 - Geometry approach
- › Use cases
 - Calibration
 - Distance estimation
- › Optional assignment
 - Smoothing, Canny, Hough transform...



AD stack



sudo apt install python3-opencv



Perception

- › We use our eyes to drive
 - Perceiving the surrounding environment and extracting information is important for autonomous driving systems
 - Monocular and stereo cameras and other sensors allow us to achieve similar functions
 - Adaptive Cruise Control
 - Vehicle, Pedestrian and Cyclist Detection
 - Automatic Emergency Braking
 - Front and Rear Vehicle Classification
 - Worldwide Traffic Sign Detection
 - Road Edge Detection
 - Lane Departure Warning
 - Lane Keep Assist
 - General Object Detection
 - Road Surface Preview
 - Free Space Detection
 - Parking Assist
 - Small Object Detection
 - Traffic Light Detection
 - Vehicle Detection at Any Angle





Perception

- › Perceiving the surrounding environment and extracting information is important for autonomous driving systems
- › Cameras are used to identify the presence and semantic attributes of the driving environment:
 1. Localization
 2. **Environmental perception**
 - › Road user's detection, classification and tracking
 - › Traffic light detection
 - › Road sign detection
 - › **Lane detection**
 - › We can use computer vision techniques or neural networks
- › **Objective (optional assignment): identify and track the position of the lines** in different images






Monocular camera



- › A monocular camera is a camera device with one lenses
- › It is designed to capture images for enhancing the safety of the driver
- › They are generally used to detect and track objects or detect lane boundaries
- › Some companies use omni directional cameras to localize the vehicle






Monocular camera



› Metrics:

- Resolution (number of pixels that create an image)
- Field of view (horizontal and vertical angle extent that is visible to the camera)
- Dynamic range


› What is the difference between the information provided by the cameras and the rest of the sensors?






Image

- › An image is a collection of **pixels** which can be represented as a **matrix**
 - contains a numerical value representing the **color/intensity**
- › **Most color and shape techniques operate over this matrix**
- › Each image is characterized by a height, width and depth
 - Depth is the number of color channel
 - › **Red, Green, Blue (RGB)**




157	153	174	168	150	152	129	151	172	161	155	156	157	153	174	168	150	152	129	151	172	161	155	156
155	182	163	74	75	62	33	17	131	210	180	154	155	182	163	74	75	62	33	17	131	210	180	154
180	180	50	14	34	6	10	33	48	106	159	181	180	180	50	14	34	6	10	33	48	106	159	181
206	109	5	124	131	111	125	204	166	15	56	180	206	109	5	124	131	111	125	204	166	15	56	180
194	64	137	251	237	239	239	228	227	87	71	201	194	64	137	251	237	239	239	228	227	87	71	201
172	155	207	233	233	214	225	239	228	96	74	206	172	155	207	233	233	214	225	239	228	96	74	206
188	88	179	209	185	215	211	158	139	75	20	169	188	88	179	209	185	215	211	158	139	75	20	169
189	87	165	84	10	168	134	11	31	62	22	148	189	87	165	84	10	168	134	11	31	62	22	148
199	168	191	193	158	227	178	143	182	105	36	190	199	168	191	193	158	227	178	143	182	105	36	190
205	174	155	252	236	231	149	178	228	43	95	234	205	174	155	252	236	231	149	178	228	43	95	234
190	216	116	148	236	187	85	150	79	38	218	241	190	216	116	148	236	187	85	150	79	38	218	241
190	224	147	108	227	210	127	102	36	101	255	224	190	224	147	108	227	210	127	102	36	101	255	224
190	214	173	66	103	143	95	50	2	109	249	215	190	214	173	66	103	143	95	50	2	109	249	215
187	196	235	75	1	81	47	0	6	217	255	211	187	196	235	75	1	81	47	0	6	217	255	211
183	202	237	145	0	0	12	104	200	138	243	236	183	202	237	145	0	0	12	104	200	138	243	236
195	206	123	207	177	121	123	200	175	13	96	218	195	206	123	207	177	121	123	200	175	13	96	218





RGB color model

- › The RGB color model is an additive color model in which red, green, and blue light are added together in various ways to reproduce a broad array of colors
- › A color in the RGB color model is described by indicating how much of each of the red, green, and blue is included. The color is expressed as an RGB triplet (r,g,b)




RGB color table

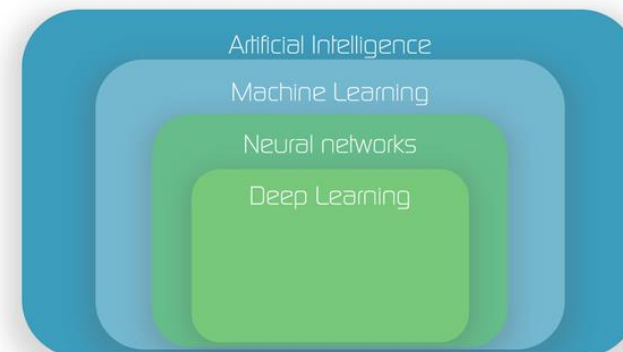

HTML / CSS Name	Hex Code #RRGGBB	Decimal Code (R,G,B)
White	#FFFFFF	(255,255,255)
Red	#FF0000	(255,0,0)
Lime	#00FF00	(0,255,0)
Blue	#0000FF	(0,0,255)



Computer vision



- › Computer vision is an interdisciplinary scientific field that deals with how computers can gain high-level understanding from digital images or videos
 - **Machine learning / neural networks**
 - › Most object detectors are implemented using neural networks/machine learning
 - Geometry/algorithmic-based approach
 - › Classic approach





Machine learning

- › Traditionally we should program by hand the instructions needed to recognize the shape of the car, the color → prone to error
- › Machine learning (ML) is the study of computer algorithms that improve automatically through **experience**
 - Detection, tracking, semantics, behavior prediction...
- › Machine learning algorithms uses a mathematical model based on sample **data (training data)**
- › To make predictions or decisions another mathematical model is applied based on the training model (**inference phase**)




```
74 75 76 75 75 76 75 74 72 70 70 70 71 72 72 71 70 71 71 68 65 74 86 88
72 74 75 75 75 76 74 73 72 70 69 69 70 73 72 72 71 71 71 67 63 62 71 76
71 74 75 75 75 76 75 73 71 68 67 68 69 72 72 71 71 70 69 66 64 63 64 66
71 73 74 75 75 76 74 70 68 65 65 66 68 72 71 70 70 69 68 65 66 67 66 65
71 72 75 75 75 76 73 70 67 65 63 63 67 71 72 70 71 68 64 63 65 68 67 64
71 72 74 74 73 73 73 70 69 76 88 95 96 90 82 75 70 66 63 62 66 68 67 64
71 72 73 72 71 70 73 77 89 105 121 123 121 119 116 104 89 68 64 64 66 67 67 64
71 71 72 70 69 70 78 103 118 120 118 115 116 118 120 120 119 62 63 66 68 67 66 64
70 71 72 70 66 69 98 123 117 108 102 100 103 104 105 110 122 113 83 69 68 65 65 63
77 76 74 67 62 78 118 119 110 108 94 87 94 99 104 108 111 116 106 76 66 64 65 62
119 116 113 104 99 109 120 116 135 135 105 92 93 104 109 109 111 117 114 85 67 64 64 61
119 141 141 137 141 113 125 111 139 137 108 104 101 105 106 105 105 112 119 111 101 98 97 93
110 116 121 120 120 127 124 102 120 121 98 109 106 104 102 99 99 104 117 111 114 116 116 115
77 78 81 79 81 102 118 102 109 116 103 108 106 106 103 94 95 106 120 116 118 118 116 118
73 73 71 70 74 89 118 111 99 107 106 101 104 110 105 93 95 113 121 87 81 84 81 81
82 82 80 79 80 82 103 130 108 95 90 93 95 98 96 93 108 122 104 69 63 66 64 64
90 91 90 88 86 80 80 112 122 106 91 90 94 94 97 110 125 112 84 66 64 65 64 64
93 94 93 91 89 83 79 81 105 125 119 114 111 109 116 120 105 86 79 74 69 66 65 65
92 91 91 90 88 86 84 82 83 91 106 116 119 119 113 96 87 88 90 89 84 78 74 72
90 91 90 89 88 86 84 84 82 79 82 86 90 93 86 82 86 92 96 98 98 96 93 87
86 87 88 87 86 85 84 84 85 85 84 83 83 84 85 86 88 92 98 99 96 98 94 94
82 84 85 84 85 84 83 83 83 82 82 82 83 83 85 86 85 88 93 92 89 92 94 92
78 81 82 83 83 82 80 80 79 79 79 79 79 80 84 84 84 87 89 88 87 88 89 90
74 75 77 79 80 79 78 78 78 78 78 78 78 80 84 83 84 86 86 86 89 89 88 88
72 72 72 74 77 79 79 79 78 78 78 78 78 79 81 83 84 82 81 87 89 90 88 88
```

Source: FF AACHEN. Alexander Ferrein
& Wikipedia



Datasets



- › Used to feed the machine learning models
- › Videos with ground truth data from different sensors
 - Generated using real-world data
 - Many options with different configurations
 - › Waymo
 - › Kitti
 - › Argoverse
 - › NuScenes
 - ›
- › They are used in many cases to validate the accuracy of research

	KITTI	NuScenes	Argo	Ours
Scenes	22	1000	113	1150
Ann. Lidar Fr.	15K	40K	22K	230K
Hours	1.5	5.5	1	6.4
3D Boxes	80K	1.4M	993k	12M
2D Boxes	80K	–	–	9.9M
Lidars	1	1	2	5
Cameras	4	6	9	5
Avg Points/Frame	120K	34K	107K	177K
LiDAR Features	1	1	1	2
Maps	No	Yes	Yes	No
Visited Area (km ²)	–	5	1.6	76

Source: Scalability in Perception for Autonomous Driving: Waymo Open Dataset. Pei Sun



Supervised learning

- › The objective is to imitate the concepts taught by the **teacher**
 - For instance let's assume that we have to teach the system to detect a car
- › We train the agent by defining where a car is in a given image
 - During the training phase, the agent learns where a car is by testing during certain intervals how good his estimation of a car is
 - A function to determine how good the estimation of the agent, is defined
- › In the inference phase the system predicts where the car is





Machine learning

- › It is very difficult to teach a system to perform as we want
 - The agent only learns what you teach it
 - › It is very easy to “cheat” an agent
- › Still a very immature technology for safety related domains





Source:
databasecultures.irmielin.org/how-to-hack-artificial-intelligence/



Camera calibration

- › Most computer vision geometric/algorithmic approaches rely in the quality of the image
 - Some cameras introduce distortion to images → this is important because of the shape of the images
 - › Radial distortion and tangential distortion (they change the shape of the image!!!)
- › To solve the distortion we have to calibrate the camera
- › The process of estimating the parameters of a camera is called camera calibration
 - **Intrinsic** are optical, geometric, and digital characteristics of the camera
 - › **Focal length**: distance between the **lens** and the image sensor
 - › **Lens distortion**: radial and tangential
 - **Extrinsic** are parameters that are external to the camera and define the **position** and **heading/orientation** of the camera with respect to the world
- › Idea: We can solve the distortion by using the intrinsic parameters of the camera. How do we obtain these?



Source: <https://de.mathworks.com/help/vision/ug/camera-calibration.html>
<http://ftp.cs.toronto.edu/pub/psala/VM/camera-parameters.pdf>



Camera calibration



$$x = PX$$

2D image point camera matrix 3D world point

$$\begin{bmatrix} f_x & 0 & c_x \\ 0 & f_y & c_y \\ 0 & 0 & 1 \end{bmatrix}$$

Intrinsic properties
(Optical Centre, scaling)

- › A camera is a mapping between the 3D world and a 2D image
- › The process of estimating the parameters of a camera is called camera calibration
 - The goal of the calibration is to find the matrix that describes the camera parameters
- › To solve this problem we provide some sample images with a well-defined pattern, i.e., **a known size and structure like a chess**
 - We find some specific points that **we already know the relative positions** (e.g. square corners in the chess board)
 - As we know the coordinates of these points in real world space and we know the coordinates in the image, so we can solve for the distortion coefficients





Camera calibration

$$x_{corrected} = x(1 + k_1r^2 + k_2r^4 + k_3r^6)$$

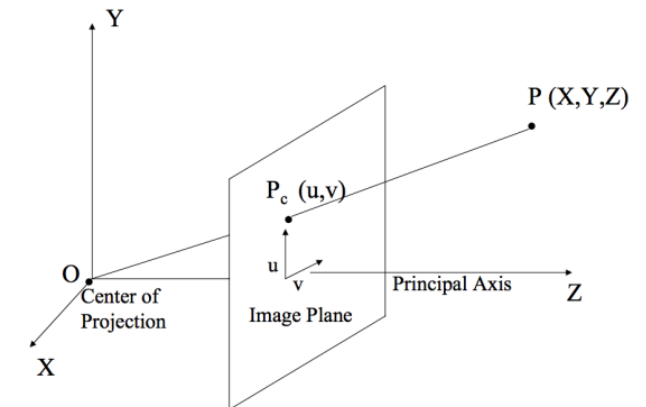
$$y_{corrected} = y(1 + k_1r^2 + k_2r^4 + k_3r^6)$$

$$x_{corrected} = x + [2p_1xy + p_2(r^2 + 2x^2)]$$

$$y_{corrected} = y + [p_1(r^2 + 2y^2) + 2p_2xy]$$

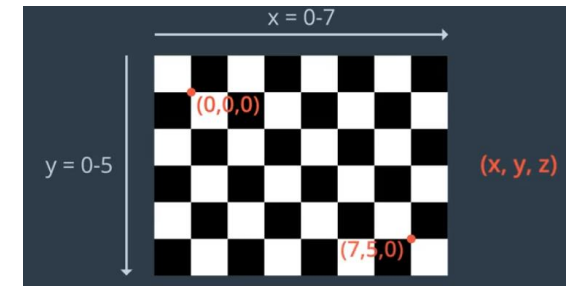
$$\text{Distortion coefficients} = (k_1 \quad k_2 \quad p_1 \quad p_2 \quad k_3)$$

- › We need some method to transform from 3D real world to 2D image coordinates
 - *A lot of math...*
 - We can extract the intrinsic parameters using the camera calibration method
- › **As we know the coordinates of the corners in the chessboard image (2D space) and where these points are in the real-world space (3D space), so we can get the values for the distortion coefficients**
 - In a nutshell: 2D image = 3D image * intrinsic parameters (you now how the 2D image is and how the 3D image is)
 - We will use the opencv function






Camera calibration



- › The process of calibrating involves two steps
 - We need a chessboard pattern in a flat surface to identify the corners
 - Try to take several pictures of the chessboard on a white surface
- › **The distortion is detected considering the difference between the shape of the images and the real size**
- › To measure how good the calibration is we can do either observe the error when projecting the points or by observe the resulting distortion



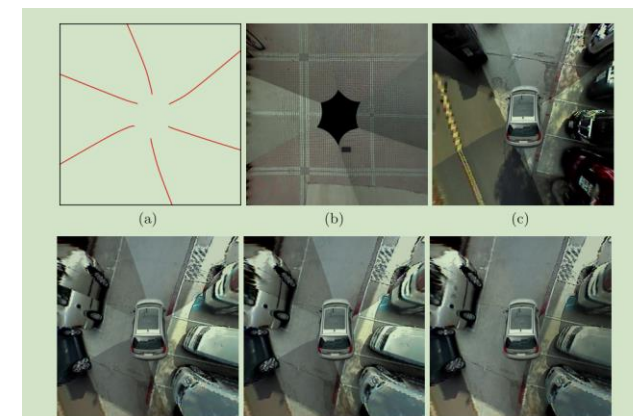
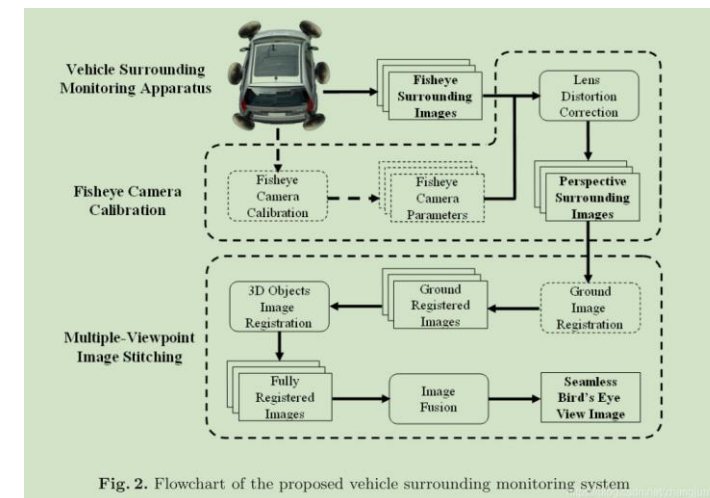

```
ret, corners = cv2.findChessboardCorners(gray, (6,8), None)
```



Bird's view

› Elon Musk recently announced a new functionality called “*Vector-space bird's eye view*”

- vision system that assists drivers by providing the panoramic image of vehicle surroundings in a bird's-eye view
- implemented stitching together 5/6 cameras surrounding the vehicle
- the perspective of the cameras are changed to the ground plane





Source: Bird's-Eye View Vision System for Vehicle Surrounding Monitoring
Liu YC., Lin KY., Chen YS. (2008)



Perspective transform

- › Perspective transform is a method to see the same scene from different viewpoint and angles
 - It is used in many applications
 - › Distance estimation
 - › Object recognition
 - › Parking assistance (bird view)
 - › Advanced lane detection
- › The farther away the object is, the smaller the objects appear
 - the **Z** information is used to transform the images
- › A warp transformation method is used





Perspective transform

1. We first choose 4 points that define the area that we would like to warp
2. We then select 4 points that define the desired rectangle plane for the warp image
3. The function **getPerspectiveTransform** returns a matrix “M” used to change the perspective
4. Next we apply the transform matrix “M” to the original image to get the warped image using the function **warpPerspective**



To Calculate the location of values of each pixel on Destination . Source Pixel values are multiplied with Transformation Matrix as given below .

Step 1 -

Transformation Matrix is multiplied with Pixel x,y coordinate values .

$$\begin{bmatrix} 2 & 0.5 & -100 \\ 0 & 2 & 0 \\ 0 & 0.005 & 1 \end{bmatrix} \cdot \begin{bmatrix} 100 \\ 100 \\ 1 \end{bmatrix} = \begin{bmatrix} 150 \\ 200 \\ 1.5 \end{bmatrix}$$

Step 2 -

Values obtained of first two rows are divided by third row to obtain (x,y) as given below.

$$\left[\frac{150}{1.5}, \frac{200}{1.5} \right] = [100, 133.33]$$

There for source to Destination Pixel location would be as Below

$$\text{Source } [100, 100] = \text{Destination } [100, 133.33]$$

Source:

<https://medium.com/analytics-vidhya/opencv-perspective-transformation-9edffefb2143>



Optional assignment

- › **Detect the lines of the road**
 - What kind of features can be helpful to identify the lines?

- Color
- Shape
- Orientation
- Position





Algorithm pipeline

1. Grayscale
2. Gaussian smoothing
3. Canny edge detection
4. Mask region of interest
5. Hough transform
6. Draw lines



Grayscale

- › The first step is to transform the image into **grayscale**
- › **Grayscale** are images with only two colors: *black and white*
 - For many applications, grayscale images are enough to perform computer vision tasks
 - › Color information doesn't help us identify important edges or other features
 - › Simplify the code
 - › Easy to compute



```
import cv2

image = cv2.imread(pathtoimage)
gray = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)


cv2.imshow('Original image', image)
cv2.imshow('Gray image', gray)

cv2.waitKey(0)
cv2.destroyAllWindows()
```




Gaussian smoothing

- › Gaussian smoothing/blurring is the result of blurring an image with a Gaussian function
- › **Used to reduce image noise or/and reduce detail**
- › In terms of image processing, any sharp edges in images are smoothed while minimizing too much blurring



```
import cv2
import numpy
src = cv2.imread('/home/img/python.png', cv2.IMREAD_UNCHANGED)
dst = cv2.GaussianBlur(src,(5,5),cv2.BORDER_DEFAULT)
cv2.imshow("Gaussian Smoothing",numpy.hstack((src, dst)))
cv2.waitKey(0) # waits until a key is pressed
cv2.destroyAllWindows() # destroys the window showing image
```



Canny Edge Detection



- › Edges correspond to a change of pixels intensity
- › The Canny edge detector is an **edge detection operator** that uses a multi-stage algorithm to detect a wide range of edges in images
 - It is mainly used to **identify the boundaries/edges** of an object in an image
 - For example, if the threshold is [0.1 0.15] then the edge pixels above the upper limit(0.15) are considered and edge pixels below the threshold(0.1) are discarded
- › The algorithm follows the next steps
 1. Noise reduction
 2. Finding Intensity Gradient of the Image
 3. Non-maximum Suppression
 4. Hysteresis Thresholding

```
edges = cv.Canny(img,low_th,high_th)
```


```
import numpy as np
import cv2 as cv
from matplotlib import pyplot as plt
img = cv.imread('messi5.png',0)
```

```
# Converting color image to grayscale image
gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
edges = cv.Canny(gray,100,200)
plt.subplot(121),plt.imshow(gray,cmap = 'gray')
plt.title('Original Image'), plt.xticks([]), plt.yticks([])
plt.subplot(122),plt.imshow(edges,cmap = 'gray')
plt.title('Edge Image'), plt.xticks([]), plt.yticks([])
plt.show()
```



Mask region of interest

- › In this phase we select the area or region that may contain the lanes
 - We can assume that the camera is fixed



```
# This time we are defining a four sided polygon to mask
imshape = image.shape
vertices = np.array([[0,imshape[0]],(460, 310), (460, 310), (imshape[1],imshape[0])], dtype=np.int32)
masked_edges=region_of_interest(edges, vertices)
```



Mask region of interest

- › We have to exclude everything that is out of the predefined area
 - “Vertices” are used to represent the region that we will retain, everything else is masked out
 - We remove everything that is out of the **region of interest**

```
def region_of_interest(img, vertices):  
    """  
    Applies an image mask.  
  
    Only keeps the region of the image defined by the polygon  
    formed from `vertices`. The rest of the image is set to black.  
    """  
  
    #defining a blank mask to start with  
    mask = np.zeros_like(img)  
  
    #defining a 3 channel or 1 channel color to fill the mask with depending on the input image  
    if len(img.shape) > 2:  
        channel_count = img.shape[2] # i.e. 3 or 4 depending on your image  
        ignore_mask_color = (255,) * channel_count  
    else:  
        ignore_mask_color = 255  
  
    #filling pixels inside the polygon defined by "vertices" with the fill color  
    cv2.fillPoly(mask, vertices, ignore_mask_color)  
  
    #returning the image only where mask pixels are nonzero  
    masked_image = cv2.bitwise_and(img, mask)  
    return masked_image
```






Hough transform (line detection)


- › How do we know that several pixels in our canny image form a line?
- › The purpose of the Hough transform technique is to find imperfect instances of objects within a certain class of shapes by a **voting procedure**
- › **Objective**: find as many lines that connect these points in image space
 - **Idea**: Collect the points that form a line and pick the best line model



image of an apple

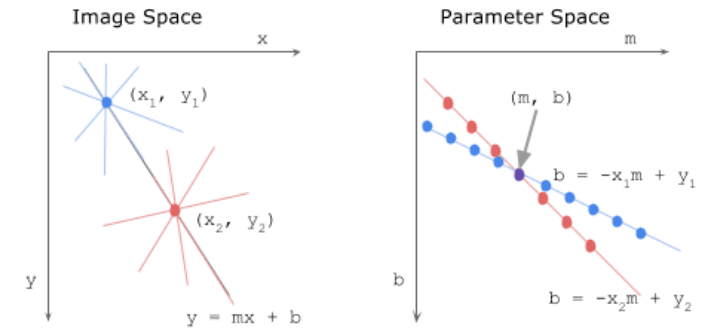


detected edges

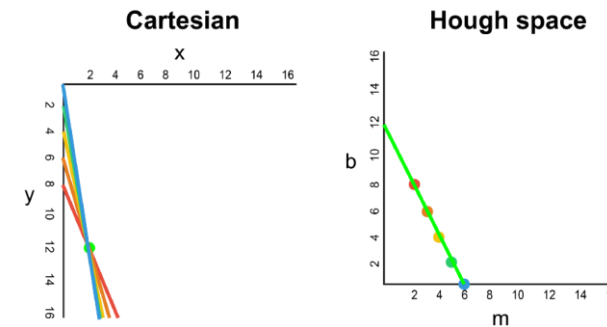





Hough transform




- › A line is a collection of points
 - EQ cartesian line: $y=mx+c$ | (**m is the slope, c is the intercept of the y axis**)
- › It is possible to use a simpler model that still characterizes a line with one point in the so-called “**m,c space**”
 - a line with the equation $y = 2x + 1$ may be represented as (2, 1) in Hough space
- › A point in the x,y space can be represented as a line in the m,c space
 - For example, a point at (2, 12) can be passed by $y = 2x + 8$, $y = 3x + 6$, $y = 4x + 4$, $y = 5x + 2$, $y = 6x$, and so on. These possible lines can be plotted in Hough space as (2, 8), (3, 6), (4, 4), (5, 2), (6, 0)





Hough transform


- › **Algorithm:** 1) iterates over each pixel, draw the line in the m, c space. 2) Count how many lines intersect. 3) The one with more votes is the best approximation of a line
- › For each white pixel, we can draw lines in the m, c space
 - The value of m tends to infinite for vertical lines (and we don't have infinite space to store infinite lines)
- › We need a different way to represent a line (**Hough space**)





Hough transform

- › (Hough space) uses the **polar coordinate system**
 - p (rho) is the perpendicular distance from the origin to the line
 - θ (theta) is the angle formed by this perpendicular line and horizontal axis
 - a line in Polar coordinate system:
 $p = x \cos \theta + y \sin \theta$
- › A line in the x, y space is still equivalent to a point in the p, θ space. **But a point in the xy space is now equivalent to a sinusoidal curve in the p, θ space**




<https://aishack.in/tutorials/hough-transform-normal/>




Hough transform

- › Image of size 30 x 30 pixels with points at (5, 25) and (20, 10)
- › The image is transformed to the Hough space by calculating ρ with a point at each angle from -90° to 90°
 - The points in Hough space make a sinusoidal curve
 - **We see that the curves in Hough space intersect at $\theta = 45^\circ$ and $\rho = 21$**
 - We can form a line in the cartesian space using as a parameters $\theta = 45^\circ$ and $\rho = 21$
- › The more curves intersect at a point, the more “votes” a line in image space will receive




point	-90°	-45°	0°	45°	90°
(5, 25)	-25	-14	5	21	25
(20, 10)	-10	7	20	21	10





Hough transform

- › You loop through every pixel of the edge image
 - If a pixel is zero, you ignore it. It's not an edge, so it can't be a line. So move on to the next pixel
 - If a pixel is nonzero, you generate its sinusoidal curve (in the p, θ space).
 - › To generate the sinusoidal curve: you take $\theta = -90$ and calculate the corresponding p value. Then, you increase the value of this cell by 1. Then you take the next θ value and calculate the next p value. And so on till $\theta = +90$. And for every such calculation, making sure they "vote"
 - The more curves intersect at a point, the more "votes" a line in image space will receive





Hough transform

- › **void cv::HoughLinesP**: Finds line segments in a binary image using the probabilistic Hough transform.

```
void cv::HoughLinesP ( InputArray image,  
                      OutputArray lines,  
                      double rho,  
                      double theta,  
                      int threshold,  
                      double minLineLength = 0 ,  
                      double maxLineGap = 0  
                      )
```

Parameters

image	8-bit, single-channel binary source image. The image may be modified by the function.
lines	Output vector of lines. Each line is represented by a 4-element vector (x_1, y_1, x_2, y_2) , where (x_1, y_1) and (x_2, y_2) are the ending points of each detected line segment.
rho	Distance resolution of the accumulator in pixels.
theta	Angle resolution of the accumulator in radians.
threshold	Accumulator threshold parameter. Only those lines are returned that get enough votes ($> \text{threshold}$).
minLineLength	Minimum line length. Line segments shorter than that are rejected.
maxLineGap	Maximum allowed gap between points on the same line to link them.

```
def hough_lines(img, rho, theta, threshold, min_line_len, max_line_gap):  
    """  
    `img` should be the output of a Canny transform.  
  
    Returns an image with hough lines drawn.  
    """  
    lines = cv2.HoughLinesP(img, rho, theta, threshold, np.array([]), minLineLength=min_line_len,  
                           maxLineGap=max_line_gap)  
    return lines
```

```
# Defining the Hough transform parameters  
# Distance resolution in pixels of the grid  
rho = 1  
# Angular resolution in radians of the grid  
theta = np.pi/180  
# Intersections in Hough grid cell  
threshold = 100  
# Minimum number of pixels making up a line  
min_line_length = 100  
# Maximum gap in pixels between connectable line segments  
max_line_gap = 10  
  
lines = hough_lines(masked_edges, rho, theta, threshold, min_line_length, max_line_gap)
```



Drawing lines



- › We need to trace a full line that connects the multiple lane markings of the image
 - **We differentiate between left and right lines**
 - › Capture negative and positive slopes (/ \)
 - › **X increases, y decreases → slope is negative → left line**
 - › **X increases, y increases → slope is positive → right line**
- › Then we interpolate all the points to form a line using the polyfit function